

STAT_ML_HW1

February 23, 2025

1 Statistical ML - Homework 1

Yonathan Amare STAT GR 5241 Spring 2025

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, cross_val_score, \
    GridSearchCV
from sklearn.linear_model import LinearRegression, Ridge, Lasso, ElasticNet, \
    ElasticNetCV, RidgeCV
from sklearn.preprocessing import StandardScaler
from sklearn.feature_selection import RFE
from sklearn.metrics import mean_squared_error
from itertools import combinations
```

1. Data Analysis

```
[2]: # Loading the data
url = "https://archive.ics.uci.edu/ml/machine-learning-databases/communities/
    \communities.data"
df = pd.read_csv(url, header=None, na_values = '?')

# Checking out the structure
df.shape, df.head()
```

```
[2]: ((1994, 128),
      \
      0      1      2      3      4      5      6      7      8      9
      \
      0      8      NaN      NaN      Lakewoodcity      1      0.19      0.33      0.02      0.90      0.12
      1      53      NaN      NaN      Tukwilacity      1      0.00      0.16      0.12      0.74      0.45
      2      24      NaN      NaN      Aberdeentown      1      0.00      0.42      0.49      0.56      0.17
      3      34      5.0      81440.0      Willingborotownship      1      0.04      0.77      1.00      0.08      0.12
      4      42      95.0      6096.0      Bethlehemtownship      1      0.01      0.55      0.02      0.95      0.09

      ...      118      119      120      121      122      123      124      125      126      127
      0      ...      0.12      0.26      0.20      0.06      0.04      0.9      0.5      0.32      0.14      0.20
```

```

1 ... 0.02 0.12 0.45 NaN NaN NaN NaN 0.00 NaN 0.67
2 ... 0.01 0.21 0.02 NaN NaN NaN NaN 0.00 NaN 0.43
3 ... 0.02 0.39 0.28 NaN NaN NaN NaN 0.00 NaN 0.12
4 ... 0.04 0.09 0.02 NaN NaN NaN NaN 0.00 NaN 0.03

```

[5 rows x 128 columns])

[3]: *# Naming the columns. Went through dataset documentation to pick which columns*
↳ to name and include in the dataframe

```

columns = [
    "state", "county", "community", "communityname", "fold",
    "population", "householdsize", "racepctblack", "racePctWhite",
    ↳ "racePctAsian",
    "racePctHisp", "agePct12t21", "agePct16t24", "agePct65up", # Removed
    ↳ agePct12t29
    "numbUrban", "pctUrban", "medIncome", "pctWWage", "pctWFarmSelf",
    "pctWInvInc", "pctWSocSec", "pctWPubAsst", "pctWRetire", "medFamInc",
    "perCapInc", "whitePerCap", "blackPerCap", "indianPerCap", "AsianPerCap",
    "HispPerCap", "NumUnderPov", "PctPopUnderPov", "PctLess9thGrade",
    ↳ "PctNotHSGrad",
    "PctBSorMore", "PctUnemployed", "PctEmploy", "PctEmplManu",
    ↳ "PctEmplProfServ",
    "PctOccupManu", "PctOccupMgmtProf", "MalePctDivorce", "MalePctNevMarr",
    ↳ "FemalePctDiv",
    "TotalPctDiv", "PersPerFam", "PctFam2Par", "PctKids2Par",
    ↳ "PctYoungKids2Par",
    "PctTeen2Par", "PctWorkMomYoungKids", "PctWorkMom", "NumIlleg", "PctIlleg",
    "NumImmig", "PctImmigRecent", "PctImmigRec5", "PctImmigRec8",
    ↳ "PctImmigRec10",
    "PctRecentImmig", "PctRecImmig5", "PctRecImmig8", "PctRecImmig10",
    ↳ "PctSpeakEnglOnly",
    "PctNotSpeakEnglWell", "PctLargHouseFam", "PctLargHouseOccup",
    ↳ "PersPerOccupHous",
    "PersPerOwnOccHous", "PersPerRentOccHous", "PctPersOwnOccup",
    ↳ "PctPersDenseHous", "PctHousLess3BR",
    "MedNumBR", "HousVacant", "PctHousOccup", "PctHousOwnOcc",
    ↳ "PctVacantBoarded",
    "PctVacMore6Mos", "MedYrHousBuilt", "PctHousNoPhone", "PctWOFullPlumb",
    ↳ "OwnOccLowQuart",
    "OwnOccMedVal", "OwnOccHiQuart", "OwnOccQrange", "RentLowQ", "RentMedian",
    "RentHighQ", "RentQrange", "MedRent", "MedRentPctHousInc",
    ↳ "MedOwnCostPctInc",
    "MedOwnCostPctIncNoMtg", "NumInShelters", "NumStreet", "PctForeignBorn",
    ↳ "PctBornSameState",
    "PctSameHouse85", "PctSameCity85", "PctSameState85", "LemasSwornFT",
    ↳ "LemasSwFTPerPop",

```

```

    "LemasSwFTFieldOps", "LemasSwFTFieldPerPop", "LemasTotalReq",
    ↪ "LemasTotReqPerPop", "PolicReqPerOffic",
    "PolicPerPop", "RacialMatchCommPol", "PctPolicWhite", "PctPolicBlack",
    ↪ "PctPolicHisp",
    "PctPolicAsian", "PctPolicMinor", "OfficAssgnDrugUnits",
    ↪ "NumKindsDrugsSeiz", "PolicAveOTWorked",
    "LandArea", "PopDens", "PctUsePubTrans", "PolicCars", "PolicOperBudg",
    "LemasPctPolicOnPatr", "LemasGangUnitDeploy", "LemasPctOfficDrugUn",
    ↪ "PolicBudgPerPop",
    "ViolentCrimesPerPop"
]

len(columns), df.shape[1]

df.columns = columns

```

```

[4]: print(f"Number of columns in df: {df.shape[1]}") # Should be 128
      print(f"Number of column names assigned: {len(columns)}") # Should be 128

```

Number of columns in df: 128
 Number of column names assigned: 128

```

[5]: # Replaced "?" with NaN and dropped columns that are entirely NaN (except
      ↪ communityname)
      communityname = df['communityname']

      df.loc[:, df.columns != "communityname"] = df.loc[:, df.columns !=
      ↪ "communityname"].replace("?", pd.NA)
      df = df.dropna(axis=1, how="all")
      df.isna().sum()

```

```

[5]: state          0
      county        1174
      community     1177
      communityname  0
      fold          0
      ...
      LemasPctPolicOnPatr  1675
      LemasGangUnitDeploy  1675
      LemasPctOfficDrugUn  0
      PolicBudgPerPop      1675
      ViolentCrimesPerPop  0
      Length: 128, dtype: int64

```

```

[6]: threshold = len(df) * 0.5 # 50% of dataset size
      df.dropna(axis = 1, thresh = threshold, inplace = True)
      df.shape[1]

```

```
[6]: 104
```

```
[7]: numeric_cols = df.columns.drop(['communityname']) # Excluding column 3
df[numeric_cols] = df[numeric_cols].apply(pd.to_numeric, errors="coerce")
df[numeric_cols] = df[numeric_cols].fillna(df[numeric_cols].mean()) # Filling
    ↪missing values with mean

df['communityname'] = communityname
```

```
[8]: df.dtypes
```

```
[8]: state                int64
communityname          object
fold                   int64
population             float64
householdsize          float64
...
LandArea              float64
PopDens              float64
PctUsePubTrans        float64
LemasPctOfficDrugUn   float64
ViolentCrimesPerPop    float64
Length: 104, dtype: object
```

```
[9]: df.isna().sum().sort_values(ascending=False).head(10)
```

```
[9]: state                0
communityname          0
PctVacMore6Mos         0
PctVacantBoarded       0
PctHousOwnOcc          0
PctHousOccup           0
HousVacant             0
MedNumBR               0
PctHousLess3BR         0
PctPersDenseHous       0
dtype: int64
```

```
[10]: df.head()
```

```
[10]:   state  communityname  fold  population  householdsize  racepctblack  \
0      8      Lakewoodcity    1         0.19           0.33         0.02
1     53      Tukwilacity    1         0.00           0.16         0.12
2     24      Aberdeentown    1         0.00           0.42         0.49
3     34  Willingborotownship    1         0.04           0.77         1.00
4     42   Bethlehemtownship    1         0.01           0.55         0.02
```

	racePctWhite	racePctAsian	racePctHisp	agePct12t21	...	PctForeignBorn	\
0	0.90	0.12	0.17	0.34	...	0.12	
1	0.74	0.45	0.07	0.26	...	0.21	
2	0.56	0.17	0.04	0.39	...	0.14	
3	0.08	0.12	0.10	0.51	...	0.19	
4	0.95	0.09	0.05	0.38	...	0.11	

	PctBornSameState	PctSameHouse85	PctSameCity85	PctSameState85	LandArea	\
0	0.42	0.50	0.51	0.64	0.12	
1	0.50	0.34	0.60	0.52	0.02	
2	0.49	0.54	0.67	0.56	0.01	
3	0.30	0.73	0.64	0.65	0.02	
4	0.72	0.64	0.61	0.53	0.04	

	PopDens	PctUsePubTrans	LemasPctOfficDrugUn	ViolentCrimesPerPop
0	0.26	0.20	0.32	0.20
1	0.12	0.45	0.00	0.67
2	0.21	0.02	0.00	0.43
3	0.39	0.28	0.00	0.12
4	0.09	0.02	0.00	0.03

[5 rows x 104 columns]

```
[11]: X = df.drop(['communityname', 'ViolentCrimesPerPop'], axis = 1) # Features
      y = df['ViolentCrimesPerPop'] # Target: violent crime rate
```

```
[12]: X.shape, y.shape
```

```
[12]: ((1994, 102), (1994,))
```

```
[13]: # 4. Split data
      X_train, X_temp, y_train, y_temp = train_test_split(X, y, test_size=0.4,
      ↪random_state=42)
      X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp, test_size=0.5,
      ↪random_state=42)

      # 5. Standardize features
      scaler = StandardScaler()
      X_train_scaled = pd.DataFrame(scaler.fit_transform(X_train), columns=X_train.
      ↪columns)
      X_val_scaled = pd.DataFrame(scaler.transform(X_val), columns=X_val.columns)
      X_test_scaled = pd.DataFrame(scaler.transform(X_test), columns=X_test.columns)

      X_train.shape, X_val.shape, X_test.shape
```

```
[13]: ((1196, 102), (399, 102), (399, 102))
```

```
[14]: non_numeric_cols = X_train.select_dtypes(include = ['object']).columns
non_numeric_cols
```

```
[14]: Index([], dtype='object')
```

```
[15]: X_train = X_train.drop(columns = non_numeric_cols)
X_val = X_val.drop(columns = non_numeric_cols)
X_test = X_test.drop(columns = non_numeric_cols)
```

(a) What are the most important features?

i. Compare and contrast the top features as determined by:

```
[16]: # OLS
ols = LinearRegression()
ols.fit(X_train, y_train)

importance = abs(ols.coef_)

feature_ranking = pd.Series(importance, index = X_train.columns).
    ↪sort_values(ascending = False)

print("Top 10 OLS Features:")
feature_ranking.head(10)
```

Top 10 OLS Features:

```
[16]: PctHousLess3BR      0.892275
PctVacMore6Mos        0.725661
blackPerCap           0.368323
pctUrban              0.359191
PersPerOwnOccHous     0.317947
MedRent               0.295365
NumImmig              0.291203
FemalePctDiv          0.281053
RentMedian            0.273867
MedNumBR              0.263495
dtype: float64
```

```
[17]: # RFE (for stepwise approach)
rfe = RFE(estimator=LinearRegression(), n_features_to_select=10)
rfe.fit(X_train, y_train)
rfe_importance = pd.Series(rfe.ranking_, index=X_train.columns)
rfe_selected = rfe_importance[rfe_importance == 1].index

print("\nFeatures selected by RFE:")
rfe_selected
```

Features selected by RFE:

```
[17]: Index(['population', 'racepctblack', 'pctUrban', 'whitePerCap', 'PctBSorMore',
          'PctUnemployed', 'PctWorkMomYoungKids', 'MedNumBR', 'PctHousOwnOcc',
          'NumStreet'],
          dtype='object')
```

```
[18]: # Lasso
from sklearn.linear_model import LassoCV

lasso = LassoCV(cv=50, random_state=42)
lasso.fit(X_train, y_train)
lasso_importance = pd.Series(abs(lasso.coef_), index=X_train.columns)
lasso_features = lasso_importance.sort_values(ascending=False)

print("Top 10 Lasso features:")
lasso_features.head(10)
```

Top 10 Lasso features:

```
[18]: PctImmigRecent      0.203366
      racepctblack      0.171315
      FemalePctDiv      0.153681
      NumStreet         0.137228
      PctWorkMomYoungKids 0.113342
      PctHousOwnOcc      0.077554
      MedNumBR          0.071715
      MedYrHousBuilt     0.064361
      PctTeen2Par        0.061865
      PctVacantBoarded   0.061485
      dtype: float64
```

```
[19]: # Elastic Net
elastic_net = ElasticNetCV(cv=5, random_state=42)
elastic_net.fit(X_train, y_train)
elastic_importance = pd.Series(abs(elastic_net.coef_), index=X_train.columns)
elastic_features = elastic_importance.sort_values(ascending=False)

print("\nTop 10 Elastic Net features:")
elastic_features.head(10)
```

Top 10 Elastic Net features:

```
[19]: PctImmigRecent      0.189986
      racepctblack      0.158529
      FemalePctDiv      0.143410
```

```

NumStreet          0.126532
PctWorkMomYoungKids 0.104400
PctHousOwnOcc       0.080582
MedYrHousBuilt      0.066296
MedNumBR            0.065641
PctVacantBoarded    0.060143
PctTeen2Par         0.056994
dtype: float64

```

```

[20]: def best_subset_selection(X, y, max_features=10):
    # 1. First, pre-screen features based on correlation with target
    correlations = abs(X.corrwith(pd.Series(y)))
    top_features = correlations.sort_values(ascending=False).head(20).index
    X_reduced = X[top_features]

    n_features = len(top_features)
    best_models = {}

    # 2. Now do best subsets on reduced feature set
    for k in range(1, min(max_features + 1, n_features + 1)):
        print(f"\nTesting {k} features...")
        best_score = float('inf')
        best_combo = None

        # Try combinations from pre-screened features
        for combo in combinations(range(n_features), k):
            X_subset = X_reduced.iloc[:, list(combo)]

            # 3. Use simple train/validation split instead of full CV
            model = LinearRegression()
            model.fit(X_subset, y)
            mse = mean_squared_error(y, model.predict(X_subset))

            if mse < best_score:
                best_score = mse
                best_combo = combo

        selected_features = [X_reduced.columns[i] for i in best_combo]
        best_models[k] = {
            'features': selected_features,
            'score': best_score
        }

        print(f"Best {k} features: {selected_features}")
        print(f"MSE: {best_score:.4f}")

    return best_models

```



```

# Run faster version
print("Starting Best Subsets Selection...")
best_models = best_subset_selection(X_train_scaled, y_train, max_features=10)

# Visualize results
plt.figure(figsize=(10, 6))
mses = [model['score'] for model in best_models.values()]
plt.plot(range(1, len(mses) + 1), mses, marker='o')
plt.xlabel('Number of Features')
plt.ylabel('MSE')
plt.title('Best Subsets Selection Results')
plt.grid(True)
plt.show()

```

Starting Best Subsets Selection...

Testing 1 features...

Best 1 features: ['PctEmplManu']

MSE: 0.0435

Testing 2 features...

Best 2 features: ['PctEmplManu', 'pctUrban']

MSE: 0.0382

Testing 3 features...

Best 3 features: ['perCapInc', 'PctEmplManu', 'pctUrban']

MSE: 0.0368

Testing 4 features...

Best 4 features: ['MedRent', 'pctWWage', 'PctEmplManu', 'pctUrban']

MSE: 0.0354

Testing 5 features...

Best 5 features: ['MedRent', 'medIncome', 'pctWWage', 'PctEmplManu', 'population']

MSE: 0.0343

Testing 6 features...

Best 6 features: ['medFamInc', 'medIncome', 'perCapInc', 'PctEmplManu', 'pctWFarmSelf', 'population']

MSE: 0.0334

Testing 7 features...

Best 7 features: ['MedRent', 'medFamInc', 'medIncome', 'perCapInc', 'PctEmplManu', 'pctWFarmSelf', 'population']

MSE: 0.0325

Testing 8 features...

Best 8 features: ['MedRent', 'medFamInc', 'medIncome', 'perCapInc', 'PctEmplManu', 'PctUsePubTrans', 'pctWFarmSelf', 'population']

MSE: 0.0321

Testing 9 features...

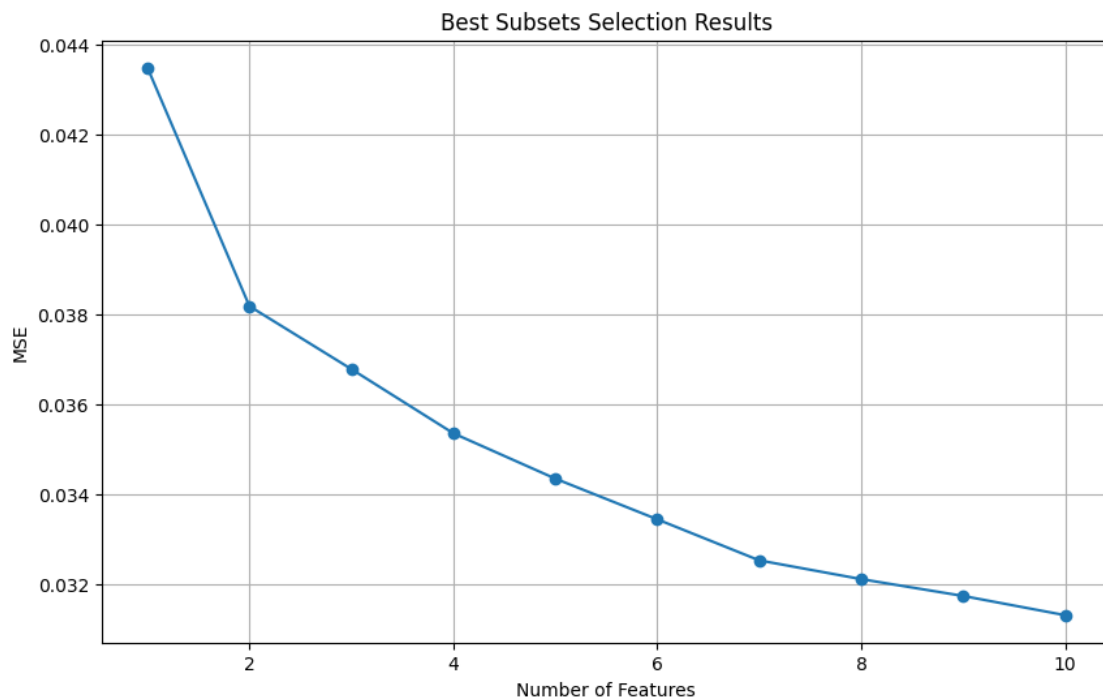
Best 9 features: ['MedRent', 'medFamInc', 'medIncome', 'perCapInc', 'PctEmplManu', 'RentMedian', 'PctUsePubTrans', 'pctWFarmSelf', 'population']

MSE: 0.0317

Testing 10 features...

Best 10 features: ['MedRent', 'medFamInc', 'medIncome', 'perCapInc', 'PctEmplManu', 'pctWPubAsst', 'RentMedian', 'PctUsePubTrans', 'pctWFarmSelf', 'population']

MSE: 0.0313



ii. Fit and visualize regularization paths for the following methods:

```
[21]: # Create range of alphas for regularization path
alphas = np.logspace(-4, 2, 100)

# Store coefficients for each path
lasso_coefs = []
elastic_coefs1 = []
```

```

elastic_coefs2 = []
ridge_coefs = []

# Calculate paths with increased max_iter and better convergence settings
for alpha in alphas:
    # Lasso with increased iterations
    lasso = Lasso(alpha=alpha, max_iter=5000, tol=1e-4)
    lasso.fit(X_train_scaled, y_train)
    lasso_coefs.append(lasso.coef_)

    # Elastic Net with l1_ratio = 0.3
    elastic1 = ElasticNet(alpha=alpha, l1_ratio=0.3, max_iter=5000, tol=1e-4)
    elastic1.fit(X_train_scaled, y_train)
    elastic_coefs1.append(elastic1.coef_)

    # Elastic Net with l1_ratio = 0.7
    elastic2 = ElasticNet(alpha=alpha, l1_ratio=0.7, max_iter=5000, tol=1e-4)
    elastic2.fit(X_train_scaled, y_train)
    elastic_coefs2.append(elastic2.coef_)

    # Ridge
    ridge = Ridge(alpha=alpha)
    ridge.fit(X_train_scaled, y_train)
    ridge_coefs.append(ridge.coef_)

# Convert to arrays
lasso_coefs = np.array(lasso_coefs)
elastic_coefs1 = np.array(elastic_coefs1)
elastic_coefs2 = np.array(elastic_coefs2)
ridge_coefs = np.array(ridge_coefs)

# Create visualization with improved styling
plt.figure(figsize=(20, 5))

# Lasso path
plt.subplot(141)
plt.semilogx(alphas, lasso_coefs)
plt.xlabel('Alpha')
plt.ylabel('Coefficients')
plt.title('Lasso Path')
plt.grid(True, alpha=0.3)

# Elastic Net path (l1_ratio=0.3)
plt.subplot(142)
plt.semilogx(alphas, elastic_coefs1)
plt.xlabel('Alpha')
plt.title('Elastic Net Path\n(l1_ratio=0.3)')


```

```
plt.grid(True, alpha=0.3)

# Elastic Net path (l1_ratio=0.7)
plt.subplot(143)
plt.semilogx(alphas, elastic_coefs2)
plt.xlabel('Alpha')
plt.title('Elastic Net Path\n(l1_ratio=0.7)')
plt.grid(True, alpha=0.3)

# Ridge path
plt.subplot(144)
plt.semilogx(alphas, ridge_coefs)
plt.xlabel('Alpha')
plt.title('Ridge Path')
plt.grid(True, alpha=0.3)

plt.tight_layout()
plt.show()
```



b. Which linear method is best for prediction?

(i) Compare the average prediction MSE on the test set for the following methods:

```
[22]: def best_subset_selection_for_prediction(X_train, y_train, X_val, y_val, X_test, y_test):
    max_features = 5 # Keep small for computational ease
    best_val_error = float('inf')
    best_test_error = None

    # Correlations
    correlations = np.array([abs(np.corrcoef(X_train[:, i], y_train)[0, 1])
                             for i in range(X_train.shape[1])])

    # Indices of top correlated features
    top_feature_indices = np.argsort(correlations)[-15:] # Get top 15 features

    for k in range(1, max_features + 1):
```

```

print(f"Testing {k} features...")
best_combo_error = float('inf')
best_combo = None

# Combinations from top features
for combo in combinations(top_feature_indices, k):
    # Fit model
    model = LinearRegression()
    model.fit(X_train[:, list(combo)], y_train)

    # Validate
    val_pred = model.predict(X_val[:, list(combo)])
    val_error = mean_squared_error(y_val, val_pred)

    if val_error < best_val_error:
        best_val_error = val_error
        test_pred = model.predict(X_test[:, list(combo)])
        best_test_error = mean_squared_error(y_test, test_pred)

return best_test_error

def evaluate_models(X, y, n_trials=10):
    results = {
        'Least Squares': [],
        'Ridge': [],
        'Lasso': [],
        'Elastic Net': [],
        'RFE': [],
        'Best Subsets': []
    }

    for trial in range(n_trials):
        print(f"Trial {trial + 1}/{n_trials}")

        # Split data
        X_temp, X_test, y_temp, y_test = train_test_split(X, y, test_size=0.2,
        ↪ random_state=trial)
        X_train, X_val, y_train, y_val = train_test_split(X_temp, y_temp,
        ↪ test_size=0.25, random_state=trial)

        # Scale features
        scaler = StandardScaler()
        X_train_scaled = scaler.fit_transform(X_train)
        X_val_scaled = scaler.transform(X_val)
        X_test_scaled = scaler.transform(X_test)

        # 1. Least Squares

```

```

lr = LinearRegression()
lr.fit(X_train_scaled, y_train)
y_pred = lr.predict(X_test_scaled)
results['Least Squares'].append(mean_squared_error(y_test, y_pred))

# 2. Ridge
ridge_cv = RidgeCV(alphas=np.logspace(-4, 4, 100))
ridge_cv.fit(X_train_scaled, y_train)
y_pred = ridge_cv.predict(X_test_scaled)
results['Ridge'].append(mean_squared_error(y_test, y_pred))

# 3. Lasso
lasso_cv = LassoCV(alphas=np.logspace(-4, 4, 100), max_iter=10000)
lasso_cv.fit(X_train_scaled, y_train)
y_pred = lasso_cv.predict(X_test_scaled)
results['Lasso'].append(mean_squared_error(y_test, y_pred))

# 4. Elastic Net
enet_cv = ElasticNetCV(alphas=np.logspace(-4, 4, 20),
                       l1_ratio=[.1, .3, .5, .7, .9],
                       max_iter=10000)
enet_cv.fit(X_train_scaled, y_train)
y_pred = enet_cv.predict(X_test_scaled)
results['Elastic Net'].append(mean_squared_error(y_test, y_pred))

# 5. RFE
rfe = RFE(estimator=LinearRegression(), n_features_to_select=10)
rfe.fit(X_train_scaled, y_train)
y_pred = rfe.predict(X_test_scaled)
results['RFE'].append(mean_squared_error(y_test, y_pred))

# 6. Best Subsets
results['Best Subsets'].append(best_subset_selection_for_prediction(
    X_train_scaled, y_train, X_val_scaled, y_val, X_test_scaled,
    y_test))

return results

# Evaluation
results = evaluate_models(X, y, n_trials=10)

# Average results
avg_results = {method: round(float(np.mean(scores)), 5) for method, scores in
    results.items()}
std_results = {method: round(float(np.std(scores)), 5) for method, scores in
    results.items()}

```

```
avg_results, std_results
```

```
Trial 1/10
Testing 1 features...
Testing 2 features...
Testing 3 features...
Testing 4 features...
Testing 5 features...
Trial 2/10
Testing 1 features...
Testing 2 features...
Testing 3 features...
Testing 4 features...
Testing 5 features...
Trial 3/10
Testing 1 features...
Testing 2 features...
Testing 3 features...
Testing 4 features...
Testing 5 features...
Trial 4/10
Testing 1 features...
Testing 2 features...
Testing 3 features...
Testing 4 features...
Testing 5 features...
Trial 5/10
Testing 1 features...
Testing 2 features...
Testing 3 features...
Testing 4 features...
Testing 5 features...
Trial 6/10
Testing 1 features...
Testing 2 features...
Testing 3 features...
Testing 4 features...
Testing 5 features...
Trial 7/10
Testing 1 features...
Testing 2 features...
Testing 3 features...
Testing 4 features...
Testing 5 features...
Trial 8/10
Testing 1 features...
Testing 2 features...
```

```

Testing 3 features...
Testing 4 features...
Testing 5 features...
Trial 9/10
Testing 1 features...
Testing 2 features...
Testing 3 features...
Testing 4 features...
Testing 5 features...
Trial 10/10
Testing 1 features...
Testing 2 features...
Testing 3 features...
Testing 4 features...
Testing 5 features...

```

```

[22]: ({'Least Squares': 0.01865,
        'Ridge': 0.01845,
        'Lasso': 0.01859,
        'Elastic Net': 0.01862,
        'RFE': 0.02033,
        'Best Subsets': 0.02116},
       {'Least Squares': 0.00182,
        'Ridge': 0.00191,
        'Lasso': 0.0019,
        'Elastic Net': 0.0019,
        'RFE': 0.00221,
        'Best Subsets': 0.00192})

```

(ii) Visualize the comparison results

```

[23]: # Create visualization
plt.figure(figsize=(12, 6))
methods = list(avg_results.keys())
means = list(avg_results.values())
stds = list(std_results.values())


plt.bar(methods, means, yerr=stds, capsize=5)
plt.xticks(rotation=45)
plt.ylabel('Mean Squared Error')
plt.title('Model Performance Comparison')
plt.tight_layout()
plt.show()

# Numerical results
for method in methods:
    print(f"{method}:")
    print(f"    Mean MSE: {avg_results[method]:.4f}")

```



```
print(f" Std MSE: {std_results[method]:.4f}")
```



Least Squares:

Mean MSE: 0.0186

Std MSE: 0.0018

Ridge:

Mean MSE: 0.0185

Std MSE: 0.0019

Lasso:

Mean MSE: 0.0186

Std MSE: 0.0019

Elastic Net:

Mean MSE: 0.0186

Std MSE: 0.0019

RFE:

Mean MSE: 0.0203

Std MSE: 0.0022

Best Subsets:

Mean MSE: 0.0212

Std MSE: 0.0019

```
[ ]:
```